A2 Written Part
By: Diego Crisafulli

Originally, I replaced every noun in the text with its seventh alphabetical neighbor, an N+7 approach, but I later realized the assignment asked for replacing only the last word of each line using GPT-2 and the P+7 technique. This shift meant focusing on a single token per line rather than every noun, and it fundamentally changed how the transformation shaped the poem. Instead of systematically changing all nouns, I discovered that focusing on the final word of each line allowed the poem's key structure to remain largely unaffected. Each line retains most of its recognizable imagery, but the final word diverges in a different way as decided by GPT-2's seventh probable token.

When experimenting with different offsets, like choosing the 10^{th} most probable token instead of the 7^{th}, I noticed that the results could be more coherent than you would expect for such a distant ranking in GPT-2's probabilities. The tenth guess may seem weird when compared to the top three, but it is still likely. GPT-2's "believability" stems from its deep comprehension of language patterns, with even lower-ranked predictions providing comprehensible outcomes. Ultimately, picking an offset like 10 offered a playful result that stayed credible while still giving the text an air of creative absurdity.

If I continued with the idea of replacing all nouns with the seventh-highest probability token, I would have needed a fully tagged text, detecting each noun, feeding the preceding context to GPT-2, and selecting the seventh token for each. To implement this, I would have separated the process into distinct steps, firstly, extract all the nouns; secondly, predict their replacements using GPT-2's seventh prediction, and third, reintegrate these replacements back into the poem. However, when I did have a different version that used all the nouns, I found the text to be completely unrecognizable, but maybe that was due to using N+7 and not P+7.