A2 Written Part
By: Diego Crisafulli

Originally, I planned to replace every noun in the text with its seventh alphabetical neighbor, an N+7 approach, but I later realized the assignment asked for replacing only the last word of each line using GPT-2 and the P+7 technique. This shift meant focusing on a single token per line rather than every noun, and it fundamentally changed how the transformation shaped the poem. Instead of systematically altering all nouns, I found that targeting just the last word of each line allows the poem's core structure to remain largely intact. Each line preserves most of its familiar imagery, yet the final word veers off in a new direction determined by GPT-2's seventh most probable token.

When experimenting with different offsets, like choosing the 10th most probable token instead of the 7th, I noticed that the results could be more coherent than you would expect for such a distant ranking in GPT-2's probabilities. The 10th guess might sound slightly odd compared to a top 3 guess, but it remains somewhat plausible. That "believability" comes from GPT-2 having a vast understanding of language patterns; even its lower-ranked predictions often give understandable results. Ultimately, picking an offset like 10 offered a playful result that stayed credible while still injecting the text with an air of creative absurdity.

If I continued with the idea of replacing all nouns with the seventh-highest probability token, I would have needed a fully tagged text, detecting each noun, feeding the preceding context to GPT-2, and selecting the seventh token for each. This would create a more thoroughly surreal effect, I would use separate it in parts, first part extract all the nouns, then predict their replacement in the seventh prediction then I would put it back in the poem.